

HOMEWORK 4

CS178

Jeremy Parnell (jimparnel)

```
In [8]: from __future__ import division
import numpy as np
import mltools as ml
import math

%matplotlib inline
import matplotlib.pyplot as plt

np.random.seed(0)
```

1 Setting up the data

```
In [9]: X = np.genfromtxt('data/X_train.txt', delimiter=None)
Y = np.genfromtxt('data/Y_train.txt', delimiter=None)
X,Y = ml.shuffleData(X,Y)
```

Problem 1.1

```
In [10]: for num in range(0,14):
          print("Feature {}:  Min: {}  Max: {}  Mean: {}  Variance: {}".format
                (num+1,np.amin(X[:,num]),np.amax(X[:,num]),np.mean(X[:,num]),np.var(X[:,
                num])))
```

Feature 1: Min: 193.5 Max: 253.0 Mean: 241.6011037 Variance: 83.4991711498

Feature 2: Min: 152.5 Max: 249.0 Mean: 227.3765713 Variance: 92.625593125

Feature 3: Min: 214.25 Max: 252.5 Mean: 241.5541505 Variance: 35.2863398033

Feature 4: Min: 152.5 Max: 252.5 Mean: 232.82676815 Variance: 97.6257317486

Feature 5: Min: 10.0 Max: 31048.0 Mean: 3089.923365 Variance: 15651513.7564

Feature 6: Min: 0.0 Max: 13630.0 Mean: 928.25902 Variance: 3081761.81695

Feature 7: Min: 0.0 Max: 9238.0 Mean: 138.09383 Variance: 443951.746446

Feature 8: Min: 0.0 Max: 125.17 Mean: 3.2485793303 Variance: 8.21948502491

Feature 9: Min: 0.87589 Max: 19.167 Mean: 6.49865290275 Variance: 6.40504819136

Feature 10: Min: 0.0 Max: 13.23 Mean: 2.09713912048 Variance: 4.36344047061

Feature 11: Min: 0.0 Max: 66.761 Mean: 4.21766040935 Variance: 4.08637188423

Feature 12: Min: 0.0 Max: 73.902 Mean: 2.69171845215 Variance: 2.19877847436

Feature 13: Min: 0.99049 Max: 975.04 Mean: 10.2715904759 Variance: 404.646245041

Feature 14: Min: -999.9 Max: 797.2 Mean: 5.7814805 Variance: 3406.52055098

Problem 1.2

```
In [11]: Xtr, Xva, Ytr, Yva = ml.splitData(X, Y)
          Xt, Yt = Xtr[:5000], Ytr[:5000] # subsample for efficiency (you can go higher)
          XtS, params = ml.rescale(Xt) # Normalize the features
          XvS, _ = ml.rescale(Xva, params) # Normalize the features
```

```
In [12]: print("X-training:")
        for num in range(0,14):
            print("Feature {}: Min: {} Max: {} Mean: {} Variance: {}".format
                  (num+1,np.amin(XtS[:,num]),np.amax(XtS[:,num]),np.mean(XtS[:,num]),np.va
                    r(XtS[:,num])))
        print("\nX-validation:")
        for num in range(0,14):
            print("Feature {}: Min: {} Max: {} Mean: {} Variance: {}".format
                  (num+1,np.amin(XvS[:,num]),np.amax(XvS[:,num]),np.mean(XvS[:,num]),np.va
                    r(XvS[:,num])))
```

X-training:

Feature 1: Min: -4.42216573151 Max: 1.24467764141 Mean: 1.0618350643
2e-14 Variance: 1.0
Feature 2: Min: -3.83799540084 Max: 1.8142505917 Mean: 8.15703060653
e-16 Variance: 1.0
Feature 3: Min: -4.59918459307 Max: 1.80668179413 Mean: -3.058886477
45e-14 Variance: 1.0
Feature 4: Min: -2.91081642991 Max: 1.95449774254 Mean: -1.167990149
04e-14 Variance: 1.0
Feature 5: Min: -0.779511378114 Max: 7.30095388843 Mean: -3.19744231
092e-17 Variance: 1.0
Feature 6: Min: -0.516235100982 Max: 7.37342139706 Mean: 7.105427357
6e-18 Variance: 1.0
Feature 7: Min: -0.20010710502 Max: 13.7671968271 Mean: -4.263256414
56e-18 Variance: 1.0
Feature 8: Min: -1.13819869133 Max: 7.35307846764 Mean: 1.5816681298
e-15 Variance: 1.0
Feature 9: Min: -2.10058928482 Max: 4.72658990213 Mean: -1.890043677
12e-15 Variance: 1.0
Feature 10: Min: -0.989591373122 Max: 5.43214474257 Mean: -1.8900436
7712e-15 Variance: 1.0
Feature 11: Min: -2.10536921641 Max: 7.41739991327 Mean: 1.364242052
66e-15 Variance: 1.0
Feature 12: Min: -1.94981434836 Max: 6.11287976902 Mean: 3.182520913
47e-15 Variance: 1.0
Feature 13: Min: -0.375997776162 Max: 37.4187664809 Mean: -2.0889956
4313e-16 Variance: 1.0
Feature 14: Min: -16.3042146041 Max: 12.7847476615 Mean: -1.31450406
116e-17 Variance: 1.0

X-validation:

Feature 1: Min: -5.16853046843 Max: 1.24467764141 Mean: -0.010605234
7626 Variance: 1.01024379009
Feature 2: Min: -3.94266662292 Max: 2.23293548003 Mean: -0.020506610
139 Variance: 1.0035701127
Feature 3: Min: -4.4782240948 Max: 1.82684187717 Mean: -0.0021482668
4922 Variance: 0.986731039437
Feature 4: Min: -2.91081642991 Max: 1.9666711725 Mean: -0.0196728968
419 Variance: 0.994919186428
Feature 5: Min: -0.779511378114 Max: 7.30095388843 Mean: 0.022502187
4675 Variance: 1.08000814085
Feature 6: Min: -0.516235100982 Max: 7.37342139706 Mean: 0.021938165
2465 Variance: 1.07093091244
Feature 7: Min: -0.20010710502 Max: 13.7671968271 Mean: 0.0063063027
3877 Variance: 1.01268888581
Feature 8: Min: -1.13819869133 Max: 8.79005323252 Mean: -0.013491928
0669 Variance: 0.983746400605
Feature 9: Min: -2.20331072131 Max: 4.63570002556 Mean: 0.0200222407
51 Variance: 0.997988030777
Feature 10: Min: -0.989591373122 Max: 5.43214474257 Mean: 0.01896850
95889 Variance: 1.01208973842
Feature 11: Min: -2.10536921641 Max: 31.1851785066 Mean: 0.001453886
91939 Variance: 1.04616651754
Feature 12: Min: -1.94981434836 Max: 51.875770174 Mean: 0.0075750868
4929 Variance: 1.173569622
Feature 13: Min: -0.378031482006 Max: 37.4187664809 Mean: -0.0176734
760563 Variance: 0.538504761236

Feature 14: Min: -16.3042146041 Max: 12.8463871288 Mean: 0.010115077
1271 Variance: 0.865691072192


2 Linear Classifiers

Problem 2.1

```
In [6]: reg = [0.0,0.1,0.2,0.3,0.4,0.5,0.6,0.7,0.8,0.9,1.0]
trAUC = np.zeros(11)
valAUC = np.zeros(11)

for i,r in enumerate(reg):
    learner = ml.linearC.linearClassify()
    learner.train(XtS, Yt, reg=r, initStep=0.5, stopTol=1e-6, stopIter=1
00)
    trAUC[i] = learner.auc(XtS, Yt)
    valAUC[i] = learner.auc(XvS, Yva)
```

```
In [7]: fig,ax=plt.subplots(1,1, figsize=(10, 8))
ax.semilogy(reg,trAUC,'r-',lw=4, marker='d', alpha=0.75, label='Training')
ax.semilogy(reg,valAUC,'g-',lw=4, marker='d', alpha=0.75, label='Test')
ax.legend(fontsize=20, loc=0)
plt.show()
```



Problem 2.2

```
In [13]: XtP = ml.transforms.fpoly(Xt,2, bias=False)
XtP, params = ml.transforms.rescale(XtP)
XvP, _ = ml.transforms.rescale( ml.transforms.fpoly(XvS,2,False), params)
```

```
In [14]: XtP.shape[1]
```

```
Out[14]: 119
```


The number of features increased to 119 because you are adding up all the possible combinations of two features out of the 14 given, while also including the possibilities that a feature could also be chosen twice when choosing the two out of the 14. For example X_1X_1 , X_2X_2 ...etc.

Problem 2.3

```
In [10]: reg = [0.0,0.1,0.2,0.3,0.4,0.5,0.6,0.7,0.8,0.9,1.0]
trAUC = np.zeros(11)
valAUC = np.zeros(11)

for i,r in enumerate(reg):
    learner = ml.linearC.linearClassify()
    learner.train(XtP, Yt, reg=r, initStep=0.01, stopTol=1e-6, stopIter=
100)
    trAUC[i] = learner.auc(XtP, Yt)
    valAUC[i] = learner.auc(XvP, Yva)
```

```
In [11]: fig,ax=plt.subplots(1,1, figsize=(8, 8))
ax.semilogy(reg,trAUC,'r-',lw=4, marker='d', alpha=0.75, label='Trainin
g')
ax.semilogy(reg,valAUC,'g-',lw=4, marker='d', alpha=0.75, label='Test')
ax.legend(fontsize=20, loc=0)
plt.show()
```




3 Nearest Neighbors

Problem 3.1

```
In [12]: kvals = [1,3,5,7,9]
trAUC = np.zeros(5)
valAUC = np.zeros(5)

for i, k in enumerate(kvals):
    learner = ml.knn.knnClassify()
    learner.train(XtS, Yt, K=k, alpha=0.0)
    trAUC[i] = learner.auc(XtS, Yt) # train AUC
    valAUC[i] = learner.auc(XvS, Yva)
```

```
In [13]: fig,ax=plt.subplots(1,1, figsize=(10, 8))
ax.semilogy(kvals,trAUC,'r-',lw=4, marker='d', alpha=0.75, label='Traini
ng')
ax.semilogy(kvals,valAUC,'g-',lw=4, marker='d', alpha=0.75, label='Test'
)
ax.legend(fontsize=20, loc=0)
plt.show()
```

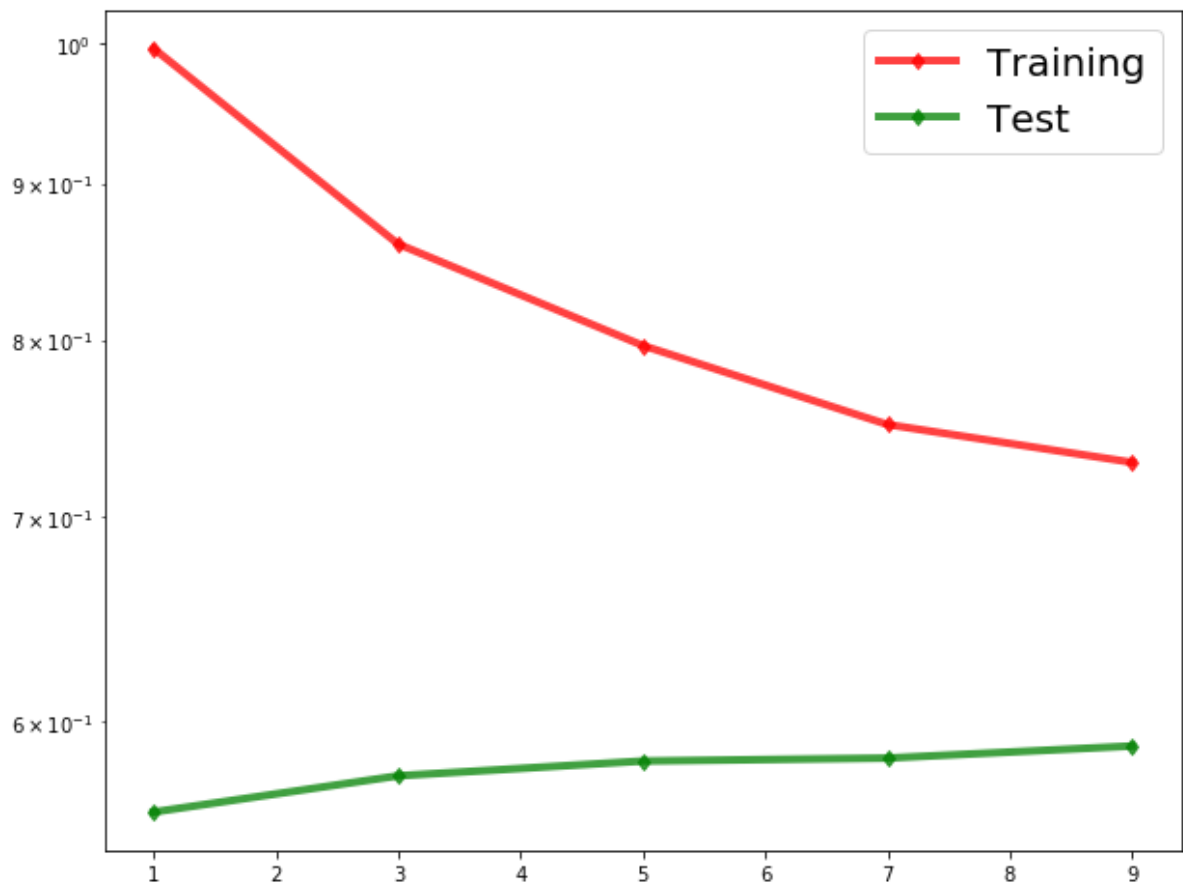


Problem 3.2


```
In [14]: kvals = [1,3,5,7,9]
trAUC = np.zeros(5)
valAUC = np.zeros(5)

for i, k in enumerate(kvals):
    learner = ml.knn.knnClassify()
    learner.train(Xt, Yt, K=k, alpha=0.0)
    trAUC[i] = learner.auc(Xt, Yt) # train AUC
    valAUC[i] = learner.auc(Xva, Yva)
```

```
In [15]: fig,ax=plt.subplots(1,1, figsize=(10, 8))
ax.semilogy(kvals,trAUC,'r-',lw=4, marker='d', alpha=0.75, label='Traini
ng')
ax.semilogy(kvals,valAUC,'g-',lw=4, marker='d', alpha=0.75, label='Test'
)
ax.legend(fontsize=20, loc=0)
plt.show()
```



Problem 3.3

```

In [16]: K = range(1,25,5) # Or something else
A = range(0,5,1) # Or something else
tr_auc = np.zeros((len(K),len(A)))
va_auc = np.zeros((len(K),len(A)))
for i,k in enumerate(K):
    for j,a in enumerate(A):
        learner = ml.knn.knnClassify()
        learner.train(XtS, Yt, K=k, alpha=a)
        tr_auc[i][j] = learner.auc(XtS, Yt) # train learner using k and
a
        va_auc[i][j] = learner.auc(XvS, Yva)

```

```

mltools\knn.py:103: RuntimeWarning: invalid value encountered in divide
prob[i,:] = count / count.sum() # save (soft) results


```

```

In [17]: # Now plot it
print("TRAINING")
f, ax = plt.subplots(1, 1, figsize=(8, 5))
cax = ax.matshow(tr_auc, interpolation='nearest')
f.colorbar(cax)
ax.set_xticklabels(['']+A)
ax.set_yticklabels(['']+K)
plt.show()


```

TRAINING



```
In [18]: print("TEST")
f, ax = plt.subplots(1, 1, figsize=(8, 5))
cax = ax.matshow(va_auc, interpolation='nearest')
f.colorbar(cax)
ax.set_xticklabels([''] + A)
ax.set_yticklabels([''] + K)
plt.show()
```

TEST



I would recommend a K value of ~21 with an alpha value of 1

4 Decision Trees

Problem 4.1

```
In [19]: mDepths = [5,15,30,45,60,70,85,100]
trAUC = np.zeros(8)
valAUC = np.zeros(8)

for i, depth in enumerate(mDepths):
    learner = ml.dtree.treeClassify(Xt, Yt, minLeaf = 1, maxDepth=depth,
    minParent = 2)
    trAUC[i] = learner.auc(XtS, Yt) # train AUC
    valAUC[i] = learner.auc(XvS, Yva)
```

```
In [20]: fig,ax=plt.subplots(1,1, figsize=(10, 8))
ax.semilogy(mDepths,trAUC,'r-',lw=4, marker='d', alpha=0.75, label='Training')
ax.semilogy(mDepths,valAUC,'g-',lw=4, marker='d', alpha=0.75, label='Test')
ax.legend(fontsize=20, loc=0)
plt.show()
```



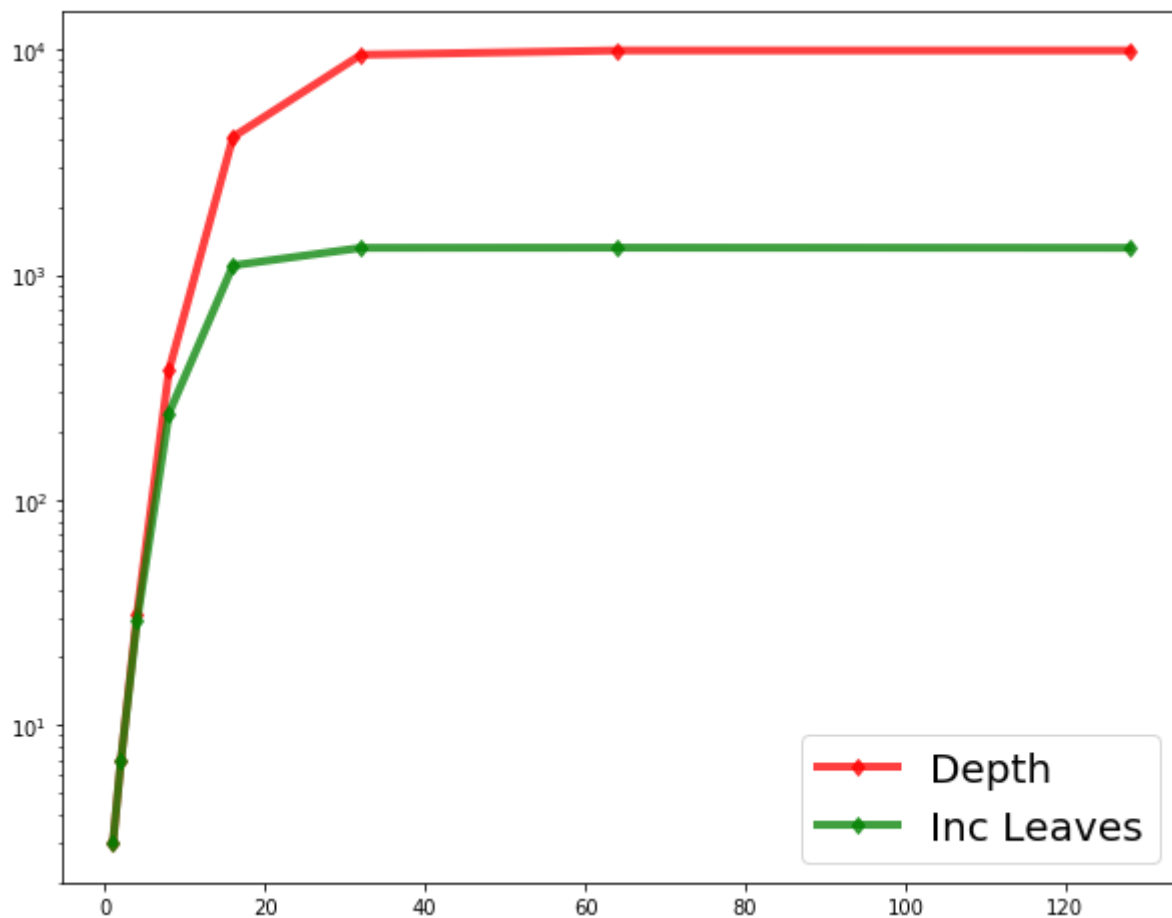
Problem 4.2

```
In [21]: mDepths = [1,2,4,8,16,32,64,128] # Or something else
num_nodes1 = np.zeros(len(mDepths))
num_nodes2 = np.zeros(len(mDepths))

for i, depth in enumerate(mDepths):
    learner = ml.dtree.treeClassify(Xt, Yt, minLeaf = 1, maxDepth = depth, minParent = 2)
    num_nodes1[i] = learner.sz

for j, depth in enumerate(mDepths):
    learner = ml.dtree.treeClassify(Xt, Yt, minLeaf = 6, maxDepth = depth, minParent = 2)
    num_nodes2[j] = learner.sz
```

```
In [22]: fig,ax=plt.subplots(1,1, figsize=(10, 8))
ax.semilogy(mDepths,num_nodes1,'r-',lw=4, marker='d', alpha=0.75, label=
'Depth')
ax.semilogy(mDepths,num_nodes2,'g-',lw=4, marker='d', alpha=0.75, label=
'Inc Leaves')
ax.legend(fontsize=20, loc=0)
plt.show()
```




Problem 4.3

```
In [23]: mLeaf = [4,8,16,32,64,128] # Or something else
mPar = [10,20,30,40,50,60] # Or something else
tr_auc = np.zeros((len(mLeaf),len(mPar)))
va_auc = np.zeros((len(mLeaf),len(mPar)))

for i, leaf in enumerate(mLeaf):
    for j, par in enumerate(mPar):
        learner = ml.dtree.treeClassify(Xt, Yt, minParent = par, maxDept
h=15, minLeaf = leaf)
        tr_auc[i][j] = learner.auc(XtS, Yt) # train learner using k and
a
        va_auc[i][j] = learner.auc(XvS, Yva)
```


```
In [24]: # Now plot it
print("TRAINING")
f, ax = plt.subplots(1, 1, figsize=(8, 5))
cax = ax.matshow(tr_auc, interpolation='nearest')
f.colorbar(cax)
ax.set_xticklabels([''] + mPar)
ax.set_yticklabels([''] + mLeaf)
plt.show()
```

TRAINING



```
In [25]: print("TEST")
f, ax = plt.subplots(1, 1, figsize=(8, 5))
cax = ax.matshow(va_auc, interpolation='nearest')
f.colorbar(cax)
ax.set_xticklabels([''] + mPar)
ax.set_yticklabels([''] + mLeaf)
plt.show()
```

TEST



According to my validation plot, it looks as though you get the best AUC with a minLeaf value of ~64 no matter the size of minParent. Therefore I would recommend a minLeaf size of 64 and a minParent value of 20

5 Neural Networks

Problem 5.1

```

In [40]: trnn_auc = np.zeros((7,2))
vann_auc = np.zeros((7,2))

nn = ml.nnet.nnetClassify()
nn.init_weights([XtS.shape[1], 6,6,6, 2], 'random', XtS, Yt) # as many 1
ayers nodes you want
nn.train(XtS, Yt, stopTol=1e-8, stepsize=.25, stopIter=300)
trnn_auc[0][0] = nn.auc(XtS,Yt)
vann_auc[0][1] = nn.auc(XvS, Yva)

nn.init_weights([XtS.shape[1], 8, 2], 'random', XtS, Yt) # as many layer
s nodes you want
nn.train(XtS, Yt, stopTol=1e-8, stepsize=.25, stopIter=300)
trnn_auc[1][0] = nn.auc(XtS,Yt)
vann_auc[1][1] = nn.auc(XvS, Yva)

nn.init_weights([XtS.shape[1], 1,1,1,1,1,1, 2], 'random', XtS, Yt) # as
many layers nodes you want
nn.train(XtS, Yt, stopTol=1e-8, stepsize=.25, stopIter=300)
trnn_auc[2][0] = nn.auc(XtS,Yt)
vann_auc[2][1] = nn.auc(XvS, Yva)

nn.init_weights([XtS.shape[1], 3,3,3,3, 2], 'random', XtS, Yt) # as many
layers nodes you want
nn.train(XtS, Yt, stopTol=1e-8, stepsize=.25, stopIter=300)
trnn_auc[3][0] = nn.auc(XtS,Yt)
vann_auc[3][1] = nn.auc(XvS, Yva)

nn.init_weights([XtS.shape[1], 9,9, 2], 'random', XtS, Yt) # as many lay
ers nodes you want
nn.train(XtS, Yt, stopTol=1e-8, stepsize=.25, stopIter=300)
trnn_auc[4][0] = nn.auc(XtS,Yt)
vann_auc[4][1] = nn.auc(XvS, Yva)

nn.init_weights([XtS.shape[1], 11,11,11,11,11, 2], 'random', XtS, Yt) #
as many layers nodes you want
nn.train(XtS, Yt, stopTol=1e-8, stepsize=.25, stopIter=300)
trnn_auc[5][0] = nn.auc(XtS,Yt)
vann_auc[5][1] = nn.auc(XvS, Yva)

nn.init_weights([XtS.shape[1], 7,7,7,7,7,7, 2], 'random', XtS, Yt) # as
many layers nodes you want
nn.train(XtS, Yt, stopTol=1e-8, stepsize=.25, stopIter=300)
trnn_auc[6][0] = nn.auc(XtS,Yt)
vann_auc[6][1] = nn.auc(XvS, Yva)

```




```
it 1 : Jsur = 0.458377048022, J01 = 0.3358
it 2 : Jsur = 0.451994865405, J01 = 0.3358
it 4 : Jsur = 0.448187339906, J01 = 0.3358
it 8 : Jsur = 0.446727377114, J01 = 0.3358
it 16 : Jsur = 0.44621736512, J01 = 0.3358
it 32 : Jsur = 0.446098371863, J01 = 0.3358
it 64 : Jsur = 0.446080002652, J01 = 0.3358
it 128 : Jsur = 0.446077293816, J01 = 0.3358
it 1 : Jsur = 0.426144848162, J01 = 0.3278
it 2 : Jsur = 0.417511875933, J01 = 0.3218
it 4 : Jsur = 0.41263595587, J01 = 0.316
it 8 : Jsur = 0.408965360036, J01 = 0.3096
it 16 : Jsur = 0.40612571949, J01 = 0.3036
it 32 : Jsur = 0.403637976118, J01 = 0.3004
it 64 : Jsur = 0.400833358079, J01 = 0.299
it 128 : Jsur = 0.398018900313, J01 = 0.3
it 256 : Jsur = 0.395218610031, J01 = 0.2974
it 1 : Jsur = 0.45837704704, J01 = 0.3358
it 2 : Jsur = 0.451994865327, J01 = 0.3358
it 4 : Jsur = 0.448187339889, J01 = 0.3358
it 8 : Jsur = 0.446727377108, J01 = 0.3358
it 16 : Jsur = 0.446217365116, J01 = 0.3358
it 32 : Jsur = 0.446098371861, J01 = 0.3358
it 64 : Jsur = 0.446080002651, J01 = 0.3358
it 128 : Jsur = 0.446077293815, J01 = 0.3358
it 1 : Jsur = 0.458377047166, J01 = 0.3358
it 2 : Jsur = 0.451994865331, J01 = 0.3358
it 4 : Jsur = 0.448187339896, J01 = 0.3358
it 8 : Jsur = 0.44672737711, J01 = 0.3358
it 16 : Jsur = 0.446217365117, J01 = 0.3358
it 32 : Jsur = 0.446098371862, J01 = 0.3358
it 64 : Jsur = 0.446080002651, J01 = 0.3358
it 128 : Jsur = 0.446077293815, J01 = 0.3358
it 1 : Jsur = 0.458377057753, J01 = 0.3358
it 2 : Jsur = 0.451994889567, J01 = 0.3358
it 4 : Jsur = 0.448187343422, J01 = 0.3358
it 8 : Jsur = 0.446727167326, J01 = 0.3358
it 16 : Jsur = 0.446201578972, J01 = 0.3358
it 32 : Jsur = 0.407790261555, J01 = 0.304
it 64 : Jsur = 0.402905135391, J01 = 0.3006
it 128 : Jsur = 0.397963506857, J01 = 0.2994
it 256 : Jsur = 0.394033352582, J01 = 0.293
it 1 : Jsur = 0.45837704959, J01 = 0.3358
it 2 : Jsur = 0.451994865522, J01 = 0.3358
it 4 : Jsur = 0.448187339933, J01 = 0.3358
it 8 : Jsur = 0.446727377124, J01 = 0.3358
it 16 : Jsur = 0.446217365123, J01 = 0.3358
it 32 : Jsur = 0.446098371865, J01 = 0.3358
it 64 : Jsur = 0.446080002651, J01 = 0.3358
it 128 : Jsur = 0.446077293815, J01 = 0.3358
it 1 : Jsur = 0.458377048934, J01 = 0.3358
it 2 : Jsur = 0.451994865509, J01 = 0.3358
it 4 : Jsur = 0.448187339911, J01 = 0.3358
it 8 : Jsur = 0.446727377115, J01 = 0.3358
it 16 : Jsur = 0.446217365118, J01 = 0.3358
it 32 : Jsur = 0.446098371863, J01 = 0.3358
```

```
it 64 : Jsur = 0.446080002651, J01 = 0.3358  
it 128 : Jsur = 0.446077293815, J01 = 0.3358
```


```
In [41]: print("TRAINING")  
f, ax = plt.subplots(1, 1, figsize=(8, 5))  
cax = ax.matshow(trnn_auc, interpolation='nearest')  
f.colorbar(cax)  
plt.show()
```

TRAINING



```
In [42]: print("TEST")  
f, ax = plt.subplots(1, 1, figsize=(8, 5))  
cax = ax.matshow(vann_auc, interpolation='nearest')  
f.colorbar(cax)  
plt.show()
```

TEST



I recommend an ideal network size to consist of 2 hidden layers of 9 nodes each layer

Problem 5.2

```
In [71]: nn = ml.nnet.nnetClassify()
sig = lambda z: np.atleast_2d(z/(1 + abs(z)))
dsig = lambda z: np.atleast_2d(1/(1 + abs(z))**2)
nn.setActivation('custom', sig, dsig)

nn.init_weights([XtS.shape[1], 9,9, 2], 'random', XtS, Yt) # as many lay
ers nodes you want
nn.train(XtS, Yt, stopTol=1e-8, stepsize=.25, stopIter=300)
print("Custom Activation Function")
print("{0:>15}: {1:.4f}".format('Train AUC',nn.auc(Xt, Yt)))
print("{0:>15}: {1:.4f}".format('Validation AUC', nn.auc(Xva, Yva)))
```

```
it 1 : Jsurr = 0.458377052549, J01 = 0.3358
it 2 : Jsurr = 0.451994885831, J01 = 0.3358
it 4 : Jsurr = 0.448187343134, J01 = 0.3358
it 8 : Jsurr = 0.446727175884, J01 = 0.3358
it 16 : Jsurr = 0.446206733134, J01 = 0.3358
it 32 : Jsurr = 0.408859100998, J01 = 0.3014
it 64 : Jsurr = 0.403928242431, J01 = 0.3012
it 128 : Jsurr = 0.398614184708, J01 = 0.3
it 256 : Jsurr = 0.393933200986, J01 = 0.2948
Custom Activation Function
      Train AUC: 0.5460
      Validation AUC: 0.5502
```

```
In [72]: nn.init_weights([XtS.shape[1], 9,9, 2], 'random', XtS, Yt) # as many lay
ers nodes you want
nn.setActivation('logistic', sig, dsig)
nn.train(XtS, Yt, stopTol=1e-8, stepsize=.25, stopIter=300)
print("Logistic Activation Function")
print("{0:>15}: {1:.4f}".format('Train AUC',nn.auc(Xt, Yt)))
print("{0:>15}: {1:.4f}".format('Validation AUC', nn.auc(Xva, Yva)))
```

```
it 1 : Jsurr = 0.437644841169, J01 = 0.3118
it 2 : Jsurr = 0.431865749811, J01 = 0.3358
it 4 : Jsurr = 0.42781539194, J01 = 0.3358
it 8 : Jsurr = 0.426119548244, J01 = 0.3358
it 16 : Jsurr = 0.425426808599, J01 = 0.3358
it 32 : Jsurr = 0.42526119897, J01 = 0.3358
it 64 : Jsurr = 0.425239892151, J01 = 0.3358
it 128 : Jsurr = 0.42516013005, J01 = 0.3358
it 256 : Jsurr = 0.425081252511, J01 = 0.3358
Logistic Activation Function
      Train AUC: 0.5347
      Validation AUC: 0.5322
```

```
In [73]: nn.init_weights([XtS.shape[1], 9,9, 2], 'random', XtS, Yt) # as many layers nodes you want
nn.setActivation('htangent', sig, dsig)
nn.train(XtS, Yt, stopTol=1e-8, stepsize=.25, stopIter=300)
print("Hyperbolic Tangent Activation Function")
print("{0:>15}: {1:.4f}".format('Train AUC',nn.auc(Xt, Yt)))
print("{0:>15}: {1:.4f}".format('Validation AUC', nn.auc(Xva, Yva)))

it 1 : Jsur = 0.458377049376, J01 = 0.3358
it 2 : Jsur = 0.451994869403, J01 = 0.3358
it 4 : Jsur = 0.448187341322, J01 = 0.3358
it 8 : Jsur = 0.446727364274, J01 = 0.3358
it 16 : Jsur = 0.446217273415, J01 = 0.3358
it 32 : Jsur = 0.446097264136, J01 = 0.3358
it 64 : Jsur = 0.411325865215, J01 = 0.307
it 128 : Jsur = 0.40315722909, J01 = 0.301
it 256 : Jsur = 0.398150458564, J01 = 0.2952
Hyperbolic Tangent Activation Function
      Train AUC: 0.5752
      Validation AUC: 0.5645
```

My custom activation function performed slightly better than the logistic activation function, but overall the hyperbolic tangent activation function performed the best and gave the highest validation AUC. The difference between all three AUC values were relatively small, however.

6 Conclusions

I believe that the **Decision Tree classifier** will perform the best with a **minLeaf value of 64**, a **minParent value of 20**, and a **maxDepth value of 15**. Decision Tree classification resulted in the highest validation AUC out of all the other classifiers, which is the reason I believe it will perform best on the complete data set.

```
In [5]: Xte = np.genfromtxt('data/X_test.txt', delimiter=None)
learner = ml.dtree.treeClassify(X, Y, minParent = 20, maxDepth=15, minLeaf = 64)
Yte = np.vstack((np.arange(Xte.shape[0]), learner.predictSoft(Xte)[: ,1])).T
np.savetxt('Y_submit.txt', Yte, '%d, %.2f', header='ID,Probl', comments='', delimiter=',')
```

After making my submission to Kaggle, my predictions came out with a score of 0.70380, advancing my team up the leaderboard 119 spots to place at 43rd. My username is jmparnel and my group's name is **The Mean Squares**.

Statement of Collaboration

All the work on this homework was done by me and me alone. Some parts were a little tricky or confusing, and so for that I looked to piazza for tips to help clear up some of the confusion in order to finish this assignment. I also met up with Sergey Kochetov and Chad Lei in person on a couple of occasions to discuss certain aspects of classifiers that I didn't understand that would help me progress through this assignment. I would also ask them questions about some of the provided code that was already given in the homework if I did not understand its functionality. However, there was no sharing of personal code, only discussion about the given code.